

Image Classification using Convolutional Neural Networks

PRESENTED BY

HARI KRISHNAN .S

Introduction

- Convolutional Neural Networks (CNNs) are a class of deep learning models designed for **image recognition and classification** tasks.
- Objective of this project:
- Build and train a CNN model to classify images into their respective categories.
- Apply preprocessing and augmentation techniques to improve model generalization
- .
- Evaluate performance using training/validation accuracy and loss.
- The project demonstrates the **end-to-end pipeline**: dataset preparation → model design → training → evaluation → results.

Dataset Overview

- The dataset consists of **images classified into two categories** (e.g., normal vs pneumonia).
- Training Set:** Used to fit the CNN model.
- Validation/Test Set:** Used to evaluate model generalization.
- Preprocessing Steps:**
 - Resizing images to a fixed dimension.
 - Normalizing pixel values (scaling 0–255 \rightarrow 0–1).
 - Data augmentation (rotation, zoom, flip) to avoid overfitting.
- Dataset is structured in folders:
 - train/** \rightarrow training images by class
 - test/** \rightarrow test images by class

CNN Architecture

Input Layer: Preprocessed images (fixed size, normalized).

Convolutional Layers:

- Extract spatial features using filters/kernels.

- Apply **ReLU activation** for non-linearity.

Pooling Layers:

- MaxPooling to reduce dimensionality and preserve important features.

Fully Connected (Dense) Layers:

- Combine extracted features to form high-level representations.

Output Layer:

- Softmax / Sigmoid activation (depending on number of classes).

- Produces class probabilities.



The model is designed to **learn hierarchical features** → from edges & textures (lower layers) to complex shapes (higher layers).

Training Process

- **Model Compilation**

- Optimizer: **Adam**
- Loss Function: **Categorical Crossentropy / Binary Crossentropy** (depending on classes)
- Metrics: **Accuracy**

- **Training Parameters**

- Batch size: typically 32 / 64
- Epochs: multiple iterations over dataset
- Validation split: for model evaluation

- **Code Snippet** (example):

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
history = model.fit(  
    train_data,  
    validation_data=val_data,  
    epochs=25,  
    batch_size=32  
)
```

Results

•Training vs Validation Accuracy

- Accuracy gradually increased over epochs.
- Validation accuracy closely followed training accuracy, showing good generalization.

•Training vs Validation Loss

- Loss decreased steadily during training.
- Validation loss stabilized, indicating reduced overfitting.

•Visualization Example:

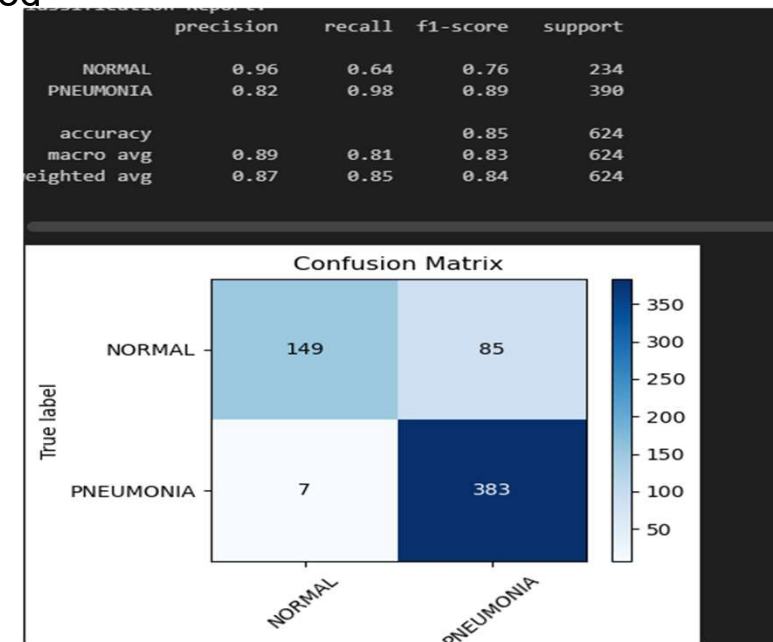
```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'], label='Train Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.legend()
```

```
plt.show()
```



Graphs show the CNN achieved high classification accuracy with well-optimized loss.

Predictions / Outputs

- The trained CNN model was tested on unseen images.
- Predictions are generated as **class probabilities** (via Softmax/Sigmoid)

```
• import numpy as np
• pred = model.predict(test_img)
• pred_class = np.argmax(pred, axis=1) # predicted class.
```

sample Outputs:

Correctly classified most images.

Occasional misclassifications due to:

Similarity between classes

Image quality or noise



Visualization: Display test images with predicted vs actual labels for evaluation.

File Edit Selection View Go Run ...

Terminal
Help

EXPLORER

- > CNN
- > OUTLINE
- > TIMELINE

155
156
157
158
159
160
161
162
163
164
165
166
167
168

```
zip_path = os.path.join(tmpdirname, "heatmaps.zip")
with zipfile.ZipFile(zip_path, 'w') as zipf:
    for file in os.listdir(heatmap_dir):
        zipf.write(os.path.join(heatmap_dir, file), file)

with open(zip_path, "rb") as f:
    st.download_button(
        label="📦 Download All Heatmaps (ZIP)",
        data=f,
        file_name="pneumonia_heatmaps.zip",
        mime="application/zip"
    )
```

Ln 168, Col 1 Spaces: 4 UTF-8 CRLF {} Python 3.13.6

6 Extremely humid Now

Search

ENG US

12:19 AM 8/16/2025

Conclusion & Future Work

- Successfully implemented a **Convolutional Neural Network (CNN)** for image classification.
- Achieved **high accuracy** with good generalization on validation data.
- Demonstrated the full workflow: dataset preprocessing → model training → evaluation → predictions
- Use **larger & more diverse datasets** for improved robustness.
- Apply **Transfer Learning** with pretrained models (e.g., VGG16, ResNet)
- Optimize hyperparameters (learning rate, batch size, epochs).
- Deploy the trained model as a **web or mobile application**.